

HR Analytics: Job Change of Data Scientists

Note:

多數特稱為類別，且部分為特徵有高比例為唯一值

特徵：

enrollee_id : 唯一ID city: 所在城市 city_development_index : 城市發展指數 (連續) gender: 性別
relevent_experience: 相關經驗 enrolled_university: 大學課程類型 education_level: 教育程度 major_discipline : 畢業學科 experience: 工作經驗年數 company_size: 當前公司員工數 company_type : 當前雇主類型
last_new_job: 上份工作與當前工作的年數差異 training_hours: 培訓時數 (連續) target: 0 不尋找新工作, 1 尋找工作

定義問題：了解特徵與 '學員願意到新公司上班而進行培訓' 之間的關係，意即具有何者特徵的學員，較有可能是為了找到新工作

In [16]:

```

from sklearn.feature_selection import RFECV
from sklearn.pipeline import Pipeline
from imblearn.under_sampling import RandomUnderSampler
from sklearn.pipeline import Pipeline
import numpy as np
from xgboost import XGBRegressor
from sklearn.model_selection import train_test_split, RandomizedSearchCV, cross_val_score
from sklearn import metrics
from sklearn.metrics import mean_squared_error, f1_score, mean_absolute_error
from sklearn import model_selection
from sklearn import feature_selection
from xgboost import XGBClassifier
from sklearn import svm, tree, linear_model, neighbors, naive_bayes, ensemble, discriminant_analysis, gaussian_process
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, QuantileTransformer, StandardScaler
from sklearn.compose import ColumnTransformer
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

```

```

train = pd.read_csv('aug_train.csv')
train.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19158 entries, 0 to 19157
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   enrollee_id                          19158 non-null  int64
 1   city                                  19158 non-null  object
 2   city_development_index                19158 non-null  float64
 3   gender                                14650 non-null  object
 4   relevent_experience                    19158 non-null  object
 5   enrolled_university                  18772 non-null  object
 6   education_level                       18698 non-null  object
 7   major_discipline                      16345 non-null  object
 8   experience                            19093 non-null  object
 9   company_size                          13220 non-null  object
10   company_type                          13018 non-null  object
11   last_new_job                          18735 non-null  object
12   training_hours                       19158 non-null  int64
13   target                                19158 non-null  float64
dtypes: float64(2), int64(2), object(10)
memory usage: 2.0+ MB


```

In [2]:


```

for col in train.columns:
    if train[col].dtypes == 'object':
        print(list(set(train[col])))
        fig, ax = plt.subplots(1, 2, figsize=(16, 10))
        train[col].value_counts().plot(kind='bar', ax=ax[0])
        plt.xticks(rotation=90)
        sns.countplot(x=col, hue='target', ax=ax[1], data=train)
        plt.show()
    else:
        continue

```



[nan, 'Part time course', 'Full time course', 'no_enrollment']



city: 過度集中 · 'city_21'較容易找新工作

experience: 大於20較多 (間隔5?) ['10', '17', nan, '14', '11', '18', '16', '4', '1', '<1', '9', '15', '12', '5', '2', '>20', '8', '7', '3', '6', '20', '13', '19']

company_size: 較分散 (小公司跳槽大公司?) ['100-500', '10/49', '1000-4999', nan, '<10', '5000-9999', '50-99', '10000+', '500-999']

education_level: Graduate較多 · 且較容易找新工作 ['High School', 'Masters', 'Phd', nan, 'Graduate', 'Primary School']

major_discipline: 過度集中STEM · 且較容易找新工作 ['Humanities', nan, 'Other', 'Arts', 'Business Degree', 'No Major', 'STEM']

company_type: 過度集中Pvt Ltd ['Pvt Ltd', 'NGO', 'Early Stage Startup', nan, 'Public Sector', 'Funded Startup', 'Other']

last_new_job: 1年較多 · never較容易找新工作 ['1', '2', nan, '3', '4', '>4', 'never']

gender: 多為男性 · 男女在target部分較無差異 ['Other', 'Male', nan, 'Female']


relevent_experience: 多具有相關經驗 · 且較不容易找新工作 ['Has relevent experience', 'No relevent experience']

In [3]:

```
train['city_development_index'].value_counts().plot(kind='bar')
# city_development_index: 和 city 分布相似
```

Out[3]:

<AxesSubplot:>

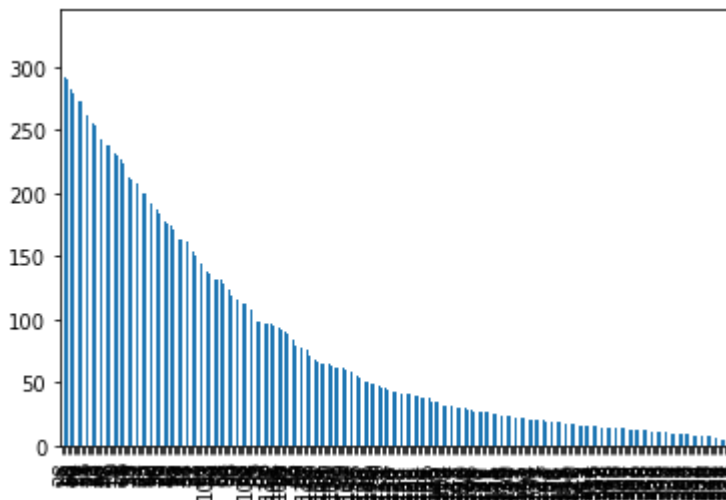


In [4]:

```
train['training_hours'].value_counts().plot(kind='bar')
# training_hours: 較分散(跟城市有關?)
```

Out[4]:

<AxesSubplot:>




In [5]:

```
train['target'].value_counts().plot(kind='bar')
# target: 0 約為1的3 倍(不尋找工作? 找不到工作?)
```

Out[5]:

<AxesSubplot:>



In [6]:

```
train.columns[train.isnull().any().values == True]
```

Out[6]:

```
Index(['gender', 'enrolled_university', 'education_level', 'major_discipline',
      'experience', 'company_size', 'company_type', 'last_new_job'],
      dtype='object')
```

In [7]:

```
pd.crosstab(train[col].fillna('missing'), train['target'].fillna('missing'), normalize=True)
```

Out[7]:

target	0.0	1.0	All
target			
0.0	0.750652	0.000000	0.750652
1.0	0.000000	0.249348	0.249348
All	0.750652	0.249348	1.000000

In [8]:

```
for col in train.columns:
    if (len(set(train[col])) < 20) & (train[col].dtype == 'object'):
        print('correlation by :', col)
        print(pd.crosstab(train[col], train['target'],
                           normalize=True, dropna=True, margins=True))
        print('-'*20, '\n')
```

```
correlation by : gender
target      0.0      1.0      All
gender
Female  0.062253  0.022253  0.084505
Male    0.696860  0.205597  0.902457
Other   0.009625  0.003413  0.013038
All     0.768737  0.231263  1.000000
-----
```

```
correlation by : relevent_experience
target      0.0      1.0      All
relevent_experience
Has relevent experience  0.565351  0.154557  0.719908
No relevent experience   0.185301  0.094791  0.280092
All                      0.750652  0.249348  1.000000
-----
```

```
correlation by : enrolled_university
target      0.0      1.0      All
enrolled_university
Full time course  0.123908  0.076231  0.200139
Part time course  0.047731  0.016088  0.063818
no_enrollment     0.580439  0.155604  0.736043
All               0.752078  0.247922  1.000000
-----
```

```
correlation by : education_level
target      0.0      1.0      All
education_level
Graduate      0.446732  0.173548  0.620280
High School   0.086801  0.021072  0.107872
Masters       0.183228  0.050005  0.233234
Phd           0.019039  0.003102  0.022141
Primary School 0.014280  0.002193  0.016472
All           0.750080  0.249920  1.000000
-----
```

```
correlation by : major_discipline
target      0.0      1.0      All
major_discipline
Arts          0.012236  0.003243  0.015479
Business Degree 0.014745  0.005262  0.020006
Humanities    0.032303  0.008626  0.040930
No Major      0.010278  0.003365  0.013643
Other         0.017069  0.006240  0.023310
STEM          0.654696  0.231936  0.886632
All           0.741328  0.258672  1.000000
-----
```

```
correlation by : company_size
target      0.0      1.0      All
company_size
10/49        0.085250  0.026021  0.111271
100-500      0.163086  0.031392  0.194478
1000-4999    0.085325  0.015129  0.100454
10000+       0.123601  0.029123  0.152723
50-99        0.191982  0.041225  0.233207
500-999      0.054841  0.011498  0.066339
5000-9999    0.034871  0.007716  0.042587
<10          0.081997  0.016944  0.098941
All          0.820953  0.179047  1.000000
-----
```

```

correlation by : company_type
target          0.0      1.0      All
company_type
Early Stage Startup  0.035413  0.010908  0.046320
Funded Startup      0.066139  0.010754  0.076894
NGO                  0.032570  0.007451  0.040022
Other                0.007067  0.002228  0.009295
Public Sector       0.057228  0.016132  0.073360
Pvt Ltd             0.617760  0.136350  0.754110
All                  0.816178  0.183822  1.000000

```

```

correlation by : last_new_job
target          0.0      1.0      All
last_new_job
1               0.315719  0.113424  0.429143
2               0.117427  0.037363  0.154790
3               0.042327  0.012330  0.054657
4               0.042754  0.012170  0.054924
>4              0.143582  0.032026  0.175607
never           0.091433  0.039445  0.130878
All             0.753243  0.246757  1.000000

```


In [9]:

```

# 資料清理
# 具有順序的編號
clean_nums = {
    'enrolled_university': {'Full time course': 3, 'Part time course': 2, 'no_enrollment': 1},
    'relevent_experience': {'Has relevent experience': 1, 'No relevent experience': 0},
    'education_level': {'Masters': 4, 'Phd': 5, 'High School': 2, 'Primary School': 1, 'Other': 0},
    'company_size': {'<10': 0, '10/49': 1, '50-99': 2, '100-500': 3, '500-999': 4, '1000-9999': 5},
    'last_new_job': {'>4': 5, 'never': 0},
    'experience': {'>20': 25, '<1': 0}
}

train = train.replace(clean_nums)

# 將city 編碼
label = LabelEncoder()
train['city'] = label.fit_transform(train['city'])

# major_discipline 與 company_type 刪除
train = train.drop(['major_discipline', 'company_type'], axis=1)

# NA 處理
# print(train['gender'].unique())
train['gender'] = train['gender'].fillna('Other')

# 假設NA代表在職中，且無經驗，教育程度為未受過教育，因此為0
train[['last_new_job', 'experience', 'education_level']] = train[['last_new_job', 'experience', 'education_level']].fillna(0)

# 假設NA代表輟學，因此同屬於'no_enrollment'
train['enrolled_university'] = train['enrolled_university'].fillna(1)

# company_size 則依據 'education_level' 進行插補
train['company_size'] = train['company_size'].fillna(train.groupby('education_level')['company_size'].transform('mean').round())

# gender 用 one-hot，照字母順序處理列名稱
onehot = OneHotEncoder()
train[['Female', 'Male', 'Other']] = onehot.fit_transform(train[['gender']]).toarray()


# training_hours 和 experience 分為5個級距
train['training_hours_cut'] = pd.qcut(train['training_hours'], 5, labels=[1, 2, 3, 4, 5])
train['experience_cut'] = pd.qcut(train['experience'].astype(int), 5, labels=[1, 2, 3, 4, 5])

for col in train.columns:
    if col not in ['gender', 'city_development_index']:
        train[col] = train[col].astype('int64')
train.head(5)

```

Out[9]:

	enrollee_id	city	city_development_index	gender	relevent_experience	enrolled_university
0	8949	5	0.920	Male	1	1
1	29725	77	0.776	Male	0	1
2	11561	64	0.624	Other	0	3
3	33241	14	0.789	Other	0	1
4	666	50	0.767	Male	1	1

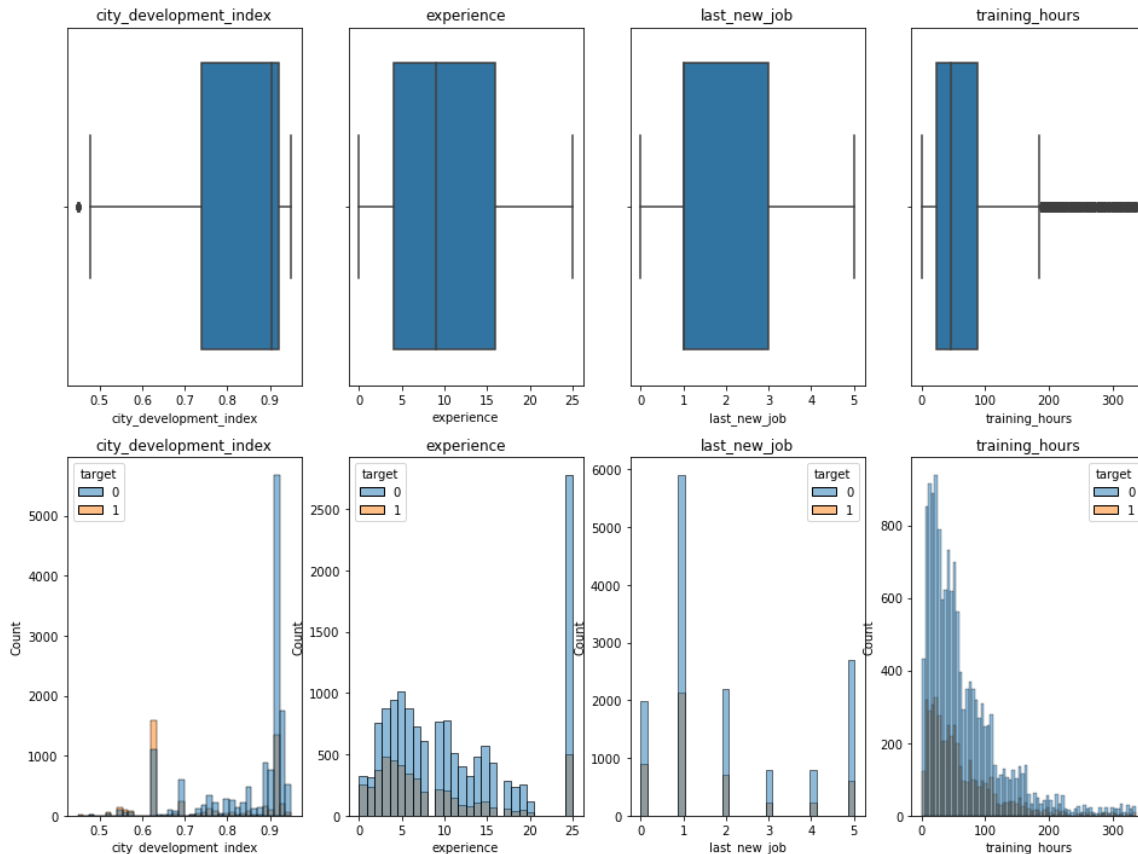


In [10]:

```
fig, ax = plt.subplots(2, 4, figsize=(16, 12))
sns.boxplot(x='city_development_index', hue='target', data=train,
            ax=ax[0, 0]).set(title='city_development_index')
sns.boxplot(x='experience', hue='target', data=train,
            ax=ax[0, 1]).set(title='experience')
sns.boxplot(x='last_new_job', hue='target', data=train,
            ax=ax[0, 2]).set(title='last_new_job')
sns.boxplot(x='training_hours', hue='target', data=train,
            ax=ax[0, 3]).set(title='training_hours')
sns.histplot(x='city_development_index', hue='target', data=train,
            ax=ax[1, 0]).set(title='city_development_index')
sns.histplot(x='experience', hue='target', data=train,
            ax=ax[1, 1]).set(title='experience')
sns.histplot(x='last_new_job', hue='target', data=train,
            ax=ax[1, 2]).set(title='last_new_job')
sns.histplot(x='training_hours', hue='target', data=train,
            ax=ax[1, 3]).set(title='training_hours')
```

Out[10]:

[Text(0.5, 1.0, 'training_hours')]



In [11]:

```

y_re = train[['target']]
x_test = train.drop('target', axis=1)

# 將資料標準化
num_col = ['experience', 'company_size', 'last_new_job', 'training_hours_cut',
           'experience_cut', 'education_level', 'enrolled_university']
skew_col = ['city_development_index', 'training_hours']

norm_num = Pipeline(
    steps=[
        ('scaler', StandardScaler())
    ]
)
skew_num = Pipeline(
    steps=[
        ('Quantile', QuantileTransformer(output_distribution='normal')),
        ('scaler', StandardScaler())
    ]
)
preprocess = ColumnTransformer(
    [
        ('norm', norm_num, num_col),
        ('skew', skew_num, skew_col)
    ]
)

X = preprocess.fit_transform(x_test)
xtest = pd.DataFrame(X, columns=num_col +
                    skew_col).join(x_test.drop(num_col+skew_col, axis=1))

x_re = xtest[['city_development_index', 'relevent_experience', 'enrolled_university', 'ec
            'experience', 'company_size', 'last_new_job', 'training_hours',
            'Female', 'Male', 'Other']]

```

In [12]:

```

# Undersampling
undersampler = RandomUnderSampler(random_state=42)
x, y = RandomUnderSampler(random_state=42).fit_resample(x_re, y_re)

x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2, random_state=432)

```

In [17]:

```

model = [
    # Ensemble Methods
    ensemble.AdaBoostClassifier(),
    ensemble.BaggingClassifier(),
    ensemble.ExtraTreesClassifier(),
    ensemble.GradientBoostingClassifier(),
    ensemble.RandomForestClassifier(),

    # GLM
    linear_model.SGDClassifier(),

    # Navies Bayes
    naive_bayes.BernoulliNB(),
    naive_bayes.GaussianNB(),

    # Nearest Neighbor
    neighbors.KNeighborsClassifier(),

    # SVM
    svm.SVC(probability=True),
    svm.LinearSVC(),

    # Trees
    tree.DecisionTreeClassifier(),
    tree.ExtraTreeClassifier(),

    # Discriminant Analysis
    discriminant_analysis.LinearDiscriminantAnalysis(),
    discriminant_analysis.QuadraticDiscriminantAnalysis(),

    # xgboost
    XGBClassifier()
]

model_result_org = pd.DataFrame(
    columns=['model', 'parameters', 'cv_accuracy', 'mae', 'mse', 'f1'])
row_index = 0

for ml in model:
    print(ml.__class__.__name__)
    ml.fit(x_train, y_train)
    y_pred = ml.predict(x_test)
    model_result_org.loc[row_index, 'model'] = ml.__class__.__name__
    model_result_org.loc[row_index, 'parameters'] = str(ml.get_params())
    model_result_org.loc[row_index, 'mae'] = mean_absolute_error(y_test, y_pred)
    model_result_org.loc[row_index, 'mse'] = mean_squared_error(y_test, y_pred)
    model_result_org.loc[row_index, 'f1'] = f1_score(y_test, y_pred)
    model_result_org.loc[row_index, 'cv_accuracy'] = cross_val_score(
        ml, x_train, y_train, cv=10, scoring='accuracy').mean()
    row_index += 1

model_result_org
# 第一次試做 : f1_score最好為 GaussianNB 之 0.513875
# 第二次試做 : f1_score最好為 GradientBoosting 為 0.751387 (部分資料標準化 + 欠採樣 + Quant
# 第三次試做 : f1_score最好為 GradientBoosting (加入cross_val做確認)

```

2	ExtraTreesClassifier	{'bootstrap': False, 'ccp_alpha': 0.0, 'class_wei...	0.686378	0.308216	0.308216	0.700559
3	GradientBoostingClassifier	{'ccp_alpha': 0.0, 'criterion': 'friedman_mse'...	0.746044	0.25798	0.25798	0.751387
4	RandomForestClassifier	{'bootstrap': True, 'ccp_alpha': 0.0, 'class_w...	0.718697	0.271586	0.271586	0.73371
5	SGDClassifier	{'alpha': 0.0001, 'average': False, 'class_wei...	0.65537	0.374673	0.374673	0.676018
6	BernoulliNB	{'alpha': 1.0, 'binarize': 0.0, 'class_prior':...	0.619126	0.379906	0.379906	0.630346
7	GaussianNB	{'priors': None, 'var_smoothing': 1e-09}	0.652624	0.349555	0.349555	0.65567

In [18]:

```
# Hyper-Parameters
```

```
model = ensemble.GradientBoostingClassifier()
model.fit(x_train, y_train)

result_cv = cross_val_score(model, x_train, y_train,
                             cv=10, scoring='accuracy').mean()

param = {'n_estimators': [40, 50, 60], 'learning_rate': [0.5, 0.1, 0.05], 'max_features':
new_model = GridSearchCV(model, param, cv=10, scoring='accuracy')
new_model.fit(x_train, y_train)

print(result_cv)
print(new_model.score(x_train, y_train))
print(new_model.best_params_)
```

```
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example usi
ng ravel().
y = column_or_1d(y, warn=True)
C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble\_gb.py:494:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example usi
ng ravel().
y = column_or_1d(y, warn=True)
C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble\_gb.py:494:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example usi
ng ravel().
y = column_or_1d(y, warn=True)
C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble\_gb.py:494:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example usi
ng ravel().
y = column_or_1d(y, warn=True)
C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble\_gb.py:494:
```

In [19]:

feature selection

```

model = ensemble.GradientBoostingClassifier()
model_rfecv = RFECV(estimator=model, step=1, cv=10)
model_rfecv.fit(x_train, y_train)
column_rfe = x_train.columns.values[model_rfecv.get_support()]

result_cv = cross_val_score(model, x_train[column_rfe], y_train,
                             cv=10, scoring='accuracy').mean()

param = {'n_estimators': [40, 50, 60], 'learning_rate': [
    0.5, 0.1, 0.05], 'max_features': ['sqrt', 'log2', 8, None], 'max_depth': [3, 6, 9]}
new_model = GridSearchCV(model, param, cv=10, scoring='accuracy')
new_model.fit(x_train[column_rfe], y_train)

```

```

print(result_cv)
print(column_rfe)
print(new_model.score(x_train[column_rfe], y_train))
print(new_model.best_params_)

```

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble_gb.py:494:

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble_gb.py:494:

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble_gb.py:494:

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble_gb.py:494:

```

['city_development_index' 'relevent_experience' 'enrolled_university' 'education_level' 'experience'
'company_size' 'last_new_job' 'training_hours' 'Male'] {'learning_rate': 0.1, 'max_depth': 3, 'max_features':
None, 'n_estimators': 50}

```

In [22]:

```

test = pd.read_csv('aug_test.csv')

test = test.replace(clean_nums)
test['city'] = label.fit_transform(test['city'])
test = test.drop(['major_discipline', 'company_type'], axis=1)
test['gender'] = test['gender'].fillna('Other')
test[['last_new_job', 'experience', 'education_level']] = test[['last_new_job', 'experience', 'education_level']].fillna(0)
test['enrolled_university'] = test['enrolled_university'].fillna(1)
test['company_size'] = test['company_size'].fillna(test.groupby('education_level')['company_size'].transform('mean').round())
test[['Female', 'Male', 'Other']] = onehot.fit_transform(test[['gender']]).toarray()
test['training_hours_cut'] = pd.qcut(test['training_hours'], 5, labels=[1, 2, 3, 4, 5])
test['experience_cut'] = pd.qcut(test['experience'].astype(int), 5, labels=[1, 2, 3, 4, 5])

for col in test.columns:
    if col not in ['gender', 'city_development_index']:
        test[col] = test[col].astype('int64')

X = preprocess.fit_transform(test)

xtest = pd.DataFrame(X, columns=num_col + skew_col).join(test.drop(num_col+skew_col, axis=1))
column_new = ['city_development_index', 'relevent_experience', 'enrolled_university', 'education_level', 'experience', 'company_size', 'last_new_job', 'training_hours_cut']
x_re = xtest[column_new]

model = ensemble.GradientBoostingClassifier(
    n_estimators=50, learning_rate=0.1, max_depth=3)
model.fit(x_train[column_new], y_train)
test["target"] = model.predict(x_re)
test.info()

```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2129 entries, 0 to 2128
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   enrollee_id                          2129 non-null   int64
1   city                                 2129 non-null   int64
2   city_development_index              2129 non-null   float64
3   gender                              2129 non-null   object
4   relevent_experience                  2129 non-null   int64
5   enrolled_university                2129 non-null   int64
6   education_level                     2129 non-null   int64
7   experience                           2129 non-null   int64
8   company_size                        2129 non-null   int64
9   last_new_job                        2129 non-null   int64
10  training_hours                      2129 non-null   int64
11  Female                              2129 non-null   int64
12  Male                                2129 non-null   int64
13  Other                               2129 non-null   int64
14  training_hours_cut                   2129 non-null   int64
15  experience_cut                       2129 non-null   int64
16  target                              2129 non-null   int64
dtypes: float64(1), int64(15), object(1)
memory usage: 282.9+ KB
```

```
C:\Users\jerry\anaconda3\lib\site-packages\sklearn\ensemble\_gb.py:494: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```